

Orchestration for efficient LLM Checkpointing within HPC

Advanced Computing Project 25/26

André Miranda PG60238

Gustavo Barros PG61527

José Soares PG60277

Miguel Carvalho PG61153

Department of Informatics – School of Engineering – University of Minho

April 24st 2026

Table of contents

- 1 Introduction
 - Context
 - Problem
 - Motivation
- 2 Approach and Architecture
 - Approach
 - Architecture
- 3 Results
- 4 Evaluation & Future Work
 - Next Steps

Context

LLM training spans several weeks, requiring periodic checkpointing to a Parallel File System (PFS) to prevent total state loss during node failures.

However, checkpointing directly to the PFS creates two problems:

- Simultaneous writes create heavy I/O contention, causing the "noisy neighbor" problem for other jobs.
- High write latency stalls training, leading to significant compute time being wasted.

Problem

Challenge	Impact
PFS Contention	I/O bottlenecks slow training as nodes compete for bandwidth.
Resilience	Hardware failures at scale risk losing weeks of training progress.
Storage Costs	Frequent PFS writes increase costs and operational overhead, which scales with larger models.

Table 1: Challenges of LLM Checkpointing to Parallel File Systems

Motivation

The goal is to design a checkpointing system that:

- Mitigates I/O contention on the PFS, ensuring fair resource sharing and Cluster Harmony.
- Provides resilience against hardware failures, minimizing training disruptions.
- Optimizes checkpointing process and storage strategies to balance performance and costs.

Approach: Policy-Driven Checkpointing

Controlled PFS Access

Mitigate I/O contention
via a centralized
Orchestrator.

Policy-Driven

Dynamic migration
based on cluster-wide
scheduling policies.

Traffic Shaping

Token Bucket to
smooth traffic and
eliminate noisy
neighbors.

Tiered Persistence

Local SSDs as a buffer,
decoupling training from
PFS latency.

Architecture




Figure 1. System Architecture.

Results

- TODO

Next Steps: Scaling and Benchmarking

Deucalion Deployment

Transition from local testing to the **Deucalion supercomputer** to validate scalability across multiple ARM nodes.

Policy Refinement

Evaluate different **scheduling policies** within the Orchestrator to find the optimal balance for PFS throughput.

Profiling

Extract performance data using the **PyTorch Profiler** during large-scale runs (already configured).

System Robustness

Ensure the Token Bucket performs reliably during massive concurrent transfers, preventing any bypass of the defined I/O policies.